

Statistical Learning and Inference Project Report

Fuming Zhang, 118033910025

Abstract—Statistical Learning

Index Terms—Statistical learning



1 INTRODUCTION

With the help of scikit-learn [1], we ...

2 DATA

2.1 Preprocessing

Scale, MinMaxScale, L1/2 Normalization

2.2 Augmentation

3 DIMENSIONALITY REDUCTION

In this section, we will introduce two dimensionality reduction methods, principle component analysis in 3.1 and factor analysis in 3.2.

3.1 Principle Component Analysis

Principle Component Analysis (abbreviation as PCA) is used to decompose a multivariate dataset in a set of successive orthogonal components that explain a maximum amount of the variance. Given a set of data in \mathbb{R}^p , denote the observations by x_1, x_2, \dots, x_N , and consider the rank- q linear model for representing them

$$f(\lambda) = \mu + V_q \lambda, \quad (1)$$

where μ is a location vector in \mathbb{R}^p , V_q is a $p \times q$ matrix with q orthogonal unit vectors as columns, and λ is a q vector of parameters. Fitting such a model to the data by least squares amounts to minimizing the reconstruction error

$$\min_{\mu, \{\lambda_i\}, V_q} \sum_{i=1}^N \|x_i - \mu - V_q \lambda_i\|^2. \quad (2)$$

We can partially optimize for μ and the λ_i to obtain

$$\hat{\mu} = \bar{x}, \quad (3)$$

$$\hat{\lambda}_i = V_q^T (x_i - \bar{x}). \quad (4)$$

This leaves us to find the orthogonal matrix V_q :

$$\min_{V_q} \sum_{i=1}^N \|(x_i - \bar{x}) - V_q V_q^T (x_i - \bar{x})\|^2. \quad (5)$$

The $p \times p$ matrix $H_q = V_q V_q^T$ is a projection matrix, and maps each point x_i onto its rank- q reconstruction $H_q x_i$, the orthogonal projection of x_i onto the subspace spanned by the columns of V_q .

3.2 Factor Analysis

The classical factor analysis model has the form

$$X = AS + \epsilon. \quad (6)$$

Here S is a vector of $q < p$ underlying latent variables or factors, A is a $p \times q$ matrix of factor loadings, and the ϵ_j are uncorrelated zero-mean disturbances. Typically the S_i and the ϵ_j are modeled as Gaussian random variables, and the model is fit by maximum likelihood. The parameters all reside in the covariance matrix

$$\Sigma = AA^T + D_\epsilon, \quad (7)$$

where $D_\epsilon = \text{diag}[\text{Var}(\epsilon_1), \dots, \text{Var}(\epsilon_p)]$. The columns of A are referred to as the factor loadings, and are used to name and interpret the factors.

Factor analysis can produce similar components (the columns of its loading matrix) to PCA. However, one can not make any general statements about these components. The main advantage for Factor Analysis over PCA is that it can model the variance in every direction of the input space independently (heteroscedastic noise). I followed the instructions given in [2] to find the best n_components of our data set in this project and the results are shown in Section ??.

4 CLASSIFICATION METHODS

In this section, we will present the classification methods we used in this project. For each method, we first review its definition and basic principles, and then introduce how we use scikit-learn [1] based implementations.

4.1 Ridge Regression

Compared with the ordinary least squares method, ridge regression shrinks the regression coefficients by imposing a penalty on their size. The ridge coefficients minimize a penalized residual sum of squares,

$$\hat{\beta}^{\text{ridge}} = \arg \min_{\beta} \left\{ \sum_{i=1}^N (y_i - \beta_0 - \sum_{j=1}^p x_{ij} \beta_j)^2 + \lambda \sum_{j=1}^p \beta_j^2 \right\}. \quad (8)$$

Here $\lambda \geq 0$ is a complexity parameter that controls the amount of shrinkage; the larger the value of λ , the greater

• E-mail: zhangfuming-alex@sytu.edu.cn

the amount of shrinkage. The coefficients are shrunk toward zero (and each other).

In this project, we use the implementation provided by 'sklearn.linear_model.RidgeClassifier'. It is a classifier using ridge regression. When dealing with a multi-class classification task, `n_class` classifiers are trained in a one-versus-all approach. Concretely, this is implemented by taking advantage of the multi-variate response support in ridge. Some important parameters of this method are introduced as follows:

- **alpha:** It is a positive float indicating the regularization strength. Regularization improves the conditioning of the problem and reduces the variance of the estimates. Large values specify stronger regularization. It is worth noting that the alpha here corresponds to C^{-1} in other linear models such as logistic regression and linear SVC.
- **normalize:** This is a boolean value which indicates whether to perform l2 normalization on the data. Its default value is false, and we have not used this parameter in our experiments. When we decide to perform data preprocessing, we will process the data before the model is trained.
- **class_weight:** The value of this parameter should be a dictionary in Python or 'balanced', which denotes the weights associated with classes. If not given, all classes are supposed to have weight one. If 'balanced' mode is used, the weights are adjusted inversely proportional to class frequencies as $\frac{n_samples}{n_classes \times np.bincount(y)}$. Since the amount of data in each class is the same in our training set, the effect of using the 'balanced' mode and the default unspecified parameter is the same.
- **solver:** The candidate values for this parameter are 'auto', 'svd', 'cholesky', 'lsqr', 'sparse_cg', 'sag' and 'saga'. It represents the method used in the computational routine. The details of these methods will not be discussed here, but we will compare the results of different methods in Section ??.

4.2 Logistic Regression

Logistic regression, despite its name, is a linear model for classification rather than regression. Logistic regression models are usually fit by maximum likelihood, using the conditional likelihood of G given X . Since $Pr(G|X)$ completely specifies the conditional distribution, the multinomial distribution is appropriate. The log-likelihood for N observations is

$$l(\theta) = \sum_{i=1}^N \log p_{g_i}(x_i; \theta), \quad (9)$$

where $p_k(x_i; \theta) = Pr(G = k|X = x_i; \theta)$. As an optimization problem, binary class L2 penalized logistic regression minimizes the following cost function:

$$\min_{\omega, c} \frac{1}{2} \beta^T \beta + C \sum_{i=1}^N \log(\exp(-y_i(\beta_i^T X_i + c)) + 1). \quad (10)$$

Similarly, L1 regularized logistic regression solves the following optimization problem:

$$\min_{\omega, c} \|\beta\|_1 + C \sum_{i=1}^N \log(\exp(-y_i(\beta_i^T X_i + c)) + 1). \quad (11)$$

In the above notation, it's assumed that the observation y_i takes values in the set $\{-1, 1\}$.

In this project, we use the implementation provided by 'sklearn.linear_model.LogisticRegression'. Table 1 and 2 summarizes the penalties supported by different solvers in logistic regression. Some important parameters of this method are introduced as follows:

- **penalty:** This parameter is used to specify the norm used in the penalization and the candidate values are 'l1' or 'l2'.
- **dual:** It is a boolean value which indicates whether to solve the principle or dual formulation problem. Dual formulation is only implemented for l2 penalty with liblinear solver. The implementer of this method recommends that we'd better set `dual=False` when `n_samples > n_features`.
- **C:** The value of this parameter must be a positive float, which represents the inverse of regularization strength. Different from the implementation of ridge regression as we have mentioned in Subsection 4.1, smaller values specify stronger regularization.
- **class_weight:** This parameter is used in the same way as described in Subsection 4.1.
- **solver:** The candidate values for this parameter are 'newton-cg', 'lbfgs', 'liblinear', 'sag' and 'saga'. It represents the algorithm used in the optimization problem. Regarding these algorithms, the following points need to be emphasized:
 - 1) For small datasets, 'liblinear' is a good choice, whereas 'sag' and 'saga' are faster for large ones.
 - 2) For multiclass problems, only 'newton-cg', 'sag', 'saga' and 'lbfgs' handle multinomial loss; 'liblinear' is limited to one-versus-rest schemes.
 - 3) 'newton-cg', 'lbfgs' and 'sag' only handle L2 penalty, whereas 'liblinear' and 'saga' handle L1 penalty.
- **multi_class:** This parameter denotes the Classification method. Candidate values are 'ovr', 'multinomial' and 'auto'. If the option chosen is 'ovr', then a binary problem is fit for each label. For 'multinomial', the loss minimized is the multinomial loss fit across the entire probability distribution.

4.3 Linear Discriminant Analysis

Linear discriminant analysis (abbreviation as LDA) is a classifier with a linear decision boundary, generated by fitting class conditional densities to the data and using Bayes' rule. The model fits a Gaussian density to each class, assuming that all classes share the same covariance matrix. The linear discriminant function has the form,

$$\delta_k(x) = x^T \Sigma^{-1} \mu_k - \frac{1}{2} \mu_k^T \Sigma^{-1} \mu_k + \log \pi_k. \quad (12)$$

TABLE 1
Penalties supported by different solvers in Logistic Regression

Penalties	liblinear	lbfgs	newton-cg	sag	saga
Multinomial + L2 penalty	no	yes	yes	yes	yes
OVR + L2 penalty	yes	yes	yes	yes	yes
Multinomial + L1 penalty	no	no	no	no	yes
OVR + L1 penalty	yes	no	no	no	yes

TABLE 2
Behaviors of different solvers in Logistic Regression

Behaviors	liblinear	lbfgs	newton-cg	sag	saga
Penalize the intercept (bad)	yes	no	no	no	no
Faster for large datasets	no	no	no	yes	yes
Robust to unscaled datasets	yes	yes	yes	no	no

The parameters of the Gaussian distributions can be estimated as follows:

- $\hat{\pi}_k = \frac{N_k}{N}$, where N_k is the number of class- k observations;
- $\hat{\delta}_k = \sum_{g_i=k} \frac{x_i}{N_k}$;
- $\hat{\Sigma} = \sum_{k=1}^K \sum_{g_i=k} \frac{(x_i - \hat{\mu}_k)(x_i - \hat{\mu}_k)^T}{N - K}$.

LDA is attractive because it has closed-form solution that can be easily computed, is inherently multiclass, has proven to work well in practice, and has no hyperparameters to tune.

In this project, we use the implementation provided by 'sklearn.discriminant_analysis.LinearDiscriminantAnalysis'. Some important parameters of this method are introduced as follows:

- **priors:** It is an array of the form (n_classes,) which represents the prior probability of each class. There are 12 categories in our data set, clearly indicating that the prior probability of each category is $\frac{1}{12}$ does not help the classification results.
- **solver:** The candidate values for this parameter are 'svd', 'lsqr' and 'eigen', which indicates the method used to solve the problem. The default solver is 'svd'. It can perform both classification and transform, and it does not rely on the calculation of the covariance matrix. This can be an advantage in situations where the number of features is large. However, the 'svd' solver cannot be used with shrinkage. The 'lsqr' solver is an efficient algorithm that only works for classification. It supports shrinkage. The 'eigen' solver is based on the optimization of the between class scatter to within class scatter ratio. It can be used for both classification and transform, and it supports shrinkage. However, the 'eigen' solver needs to compute the covariance matrix, so it might not be suitable for situations with a high number of features. We will discuss the performance of these methods in Section ??.

4.4 Support Vector Machine

Support Vector Machine (abbreviation as SVM) constructs a hyperplane or set of hyperplanes in a high- or infinite-dimensional space, which can be used for classification

and regression. Given slack variables ξ , the formulation of "standard" support vector classifier is:

$$\min_{\beta, \beta_0} \frac{1}{2} \|\beta\|^2 + C \sum_{i=1}^N \xi_i \quad (13)$$

$$\text{subject to } \xi_i \geq 0, y_i(x_i^T \beta + \beta_0) \geq M(1 - \xi_i), \forall i.$$

The Lagrange (primal) function is

$$L_P = \frac{1}{2} \|\beta\|^2 + C \sum_{i=1}^N \xi_i - \sum_{i=1}^N \mu_i \xi_i - \sum_{i=1}^N \alpha_i [y_i(x_i^T \beta + \beta_0) - (1 - \xi_i)], \quad (14)$$

which we minimize w.r.t β, β_0 and ξ_i . Setting the respective derivatives to zero, we get

$$\beta = \sum_{i=1}^N \alpha_i y_i x_i, \quad (15)$$

$$0 = \sum_{i=1}^N \alpha_i y_i, \quad (16)$$

$$\alpha_i = C - \mu_i, \forall i, \quad (17)$$

as well as the positivity constraints $\alpha_i, \mu_i, \xi_i \geq 0 \forall i$. By substituting Eq.(15)-(17) into Eq.(14), we obtain the Lagrangian dual objective function

$$L_D = \sum_{i=1}^N \alpha_i - \frac{1}{2} \sum_{i=1}^N \sum_{i'=1}^N \alpha_i \alpha_{i'} y_i y_{i'} x_i^T x_{i'}, \quad (18)$$

which gives a lower bound on the objective function for any feasible point. We maximize L_D subject to $0 \leq \alpha_i \leq C$ and $\sum_{i=1}^N \alpha_i y_i = 0$. In addition to Eq.(15)-(17), the Karush-Kuhn-Tucker conditions include the constraints

$$\alpha_i [y_i(x_i^T \beta + \beta_0) - (1 - \xi_i)] = 0, \quad (19)$$

$$\mu_i \xi_i = 0, \quad (20)$$

$$y_i(x_i^T \beta + \beta_0) - (1 - \xi_i) \geq 0, \quad (21)$$

for $i = 1, \dots, N$.

In this project, we use the implementation provided by 'sklearn.svm.LinearSVC'. Similar to SVC with parameter

kernel='linear', but implemented in terms of liblinear rather than libsvm, so it has more flexibility in the choice of penalties and loss functions and should scale better to large numbers of samples. Some important parameters of this method are introduced as follows:

- **penalty:** This parameter is used to specify the norm used in the penalization and the candidate values are 'l1' or 'l2'.
- **loss:** This parameter specifies the loss function. 'hinge' is the standard SVM loss while 'squared_hinge' is the square of the hinge loss.
- **dual:** It specifies whether to solve the dual or primal optimization problem. Prefer dual=False when $n_samples > n_features$.
- **C:** This is the penalty parameter of the error term.

4.5 K Nearest Neighbour

neighbours-based classification is a type of instance-based learning or non-generalizing learning: it does not attempt to construct a general internal model, but simply stores instances of the training data. Classification is computed from a simple majority vote of the nearest neighbours of each point: a query point is assigned the data class which has the most representatives within the nearest neighbours of the point.

The k-neighbours classification is the most commonly used technique. The optimal choice of the value k is highly data-dependent: in general a larger k suppresses the effects of noise, but makes the classification boundaries less distinct. Specifically, the k-nearest neighbour fit has the follow formulation:

$$\hat{f}(x_0) = \frac{1}{k} \sum_{l=1}^k f(x_{(l)}), \quad (22)$$

where the subscripts in parantheses (l) indicate the sequence of nearest neighbours to x_0 .

In this project, we use the implementation provided by 'sklearn.neighbors.KNeighborsClassifier'. Some important parameters of this method are introduced as follows:

- **n_neighbours:** This parameter specifies the number of neighbours to use by default for k-neighbours query.
- **weights:** It represents the weight function used in prediction. Two candidate values are 'uniform' and 'distance'. For 'uniform', all points in each neighborhood are weighted equally. For 'distance', points are weighted by inverse of their distance. In this case, closer neighbors of a query point will have a greater influence than neighbors which are further away.
- **p:** It indicates the power parameter for the Minkowski metric. Psually p is equal to 1 or 2.
- **algorithm:** This parameter denotes the algorithm used to compute the nearest neighbours. Candidate algorithms are 'ball_tree', 'kd_tree' and 'brute'.
- **leaf_size:** This parameter is passed to BallTree or KDTree, which can affect the speed of the construction and query, as well as the memory required to store the tree. The optimal value depends on the nature of the problem.

5 MODEL SELECTION

Cross Validation, Grid Search

6 EVALUATION

7 DISCUSSION

8 CONCLUSION

In this paper [3], we have proposed...

ACKNOWLEDGMENTS

The author would like to thank...

REFERENCES

- [1] F. Pedregosa, G. Varoquaux, A. Gramfort, V. Michel, B. Thirion, O. Grisel, M. Blondel, P. Prettenhofer, R. Weiss, V. Dubourg, J. Vanderplas, A. Passos, D. Cournapeau, M. Brucher, M. Perrot, and E. Duchesnay, "Scikit-learn: Machine learning in python," *Journal of Machine Learning Research*, vol. 12, pp. 2825–2830, 2011.
- [2] A. Gramfort and D. Engemann, "Model selection with probabilistic pca and factor analysis," https://scikit-learn.org/stable/auto_examples/decomposition/plot_pca_vs_fa_model_selection.html#sphx-glr-auto-examples-decomposition-plot-pca-vs-fa-model-selection-py, accessed January 18, 2019.
- [3] F. Zhang, Z. Tang, M. Chen, X. Zhou, and W. Jia, "A dynamic resource overbooking mechanism in fog computing," in *15th IEEE International Conference on Mobile Ad Hoc and Sensor Systems (MASS)*, 2018, pp. 89–97.